

Bayesian Learning (732A73) Lab 1

Hoda Fakharzadehjahromy (hodfa840), Ravinder Atia(ravat601)

Assisgnment 1

Daniel Bernoulli

Bernoulli ...again.

Let $y_1, \dots, y_n | \theta \sim \text{Bern}(\theta)$ and assume that you have obtained a sample with $s = 8$ successes in $n = 24$ trials. Assume a $\text{Beta}(\alpha_0, \beta_0)$ prior for θ and let $\alpha_0 = \beta_0 = 3$

(a) Draw random numbers from the posterior $\theta | y \sim \text{Beta}(\alpha_0 + s, \beta_0 + f)$, where $y = (y_1, \dots, y_n)$; and verify graphically that the posterior mean and standard deviation converges to the true values as the number of random draws grows large.

(b) Use simulation ($n\text{Draws} = 10000$) to compute the posterior probability $\Pr(\theta < 0.4 | y)$ and compare with the exact value.

(c) Compute the posterior distribution of the log-odds $\phi = \log \frac{\theta}{1-\theta}$ ($n\text{Draws} = 10000$)

(a)

True mean for $\text{Beta}(\alpha, \beta)$ is:

$$E[\theta] = \frac{\alpha}{\alpha + \beta}$$

True σ^2 for $\text{Beta}(\alpha, \beta)$ is:

$$\sigma^2(\theta) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$$

The prior function for this problem is:

$$p(\theta) \propto \theta^{\alpha-1}(1 - \theta)^{\beta-1}$$

The likelihood function is :

$$p(y_1, \dots, y_n | \theta) \propto \prod_{i=1}^n p(y_i | \theta) = \prod_{i=1}^n \theta^{y_i} (1 - \theta)^{1-y_i} = \theta^s (1 - \theta)^f$$

posterior is :

posterior

$$p(\theta | y) \propto \text{Beta}(\alpha_0 + s, \beta_0 + f) = \text{Beta}(11, 19)$$

```
In [1]: s = 8
n = 24
f = 16
alpha_0 = beta_0 = 3
```

```
In [2]: NDraws = 10000
res = matrix(, NDraws, 3)
colnames(res) = c("iter", "mean", "sd")
alpha_n = alpha_0 + s
beta_n = beta_0 + f
TrueMean = (alpha_n) / (alpha_n + beta_n)
TrueVar = sqrt((alpha_n * beta_n) / ((alpha_n + beta_n)^2 * (alpha_n + beta_n + 1)))
TrueMean
TrueVar
```

0.366666666666667
0.0865507910219387

```
In [3]: PostB <- function(alpha = alpha_n, beta = beta_n, n = NDraws) {
  set.seed(43)

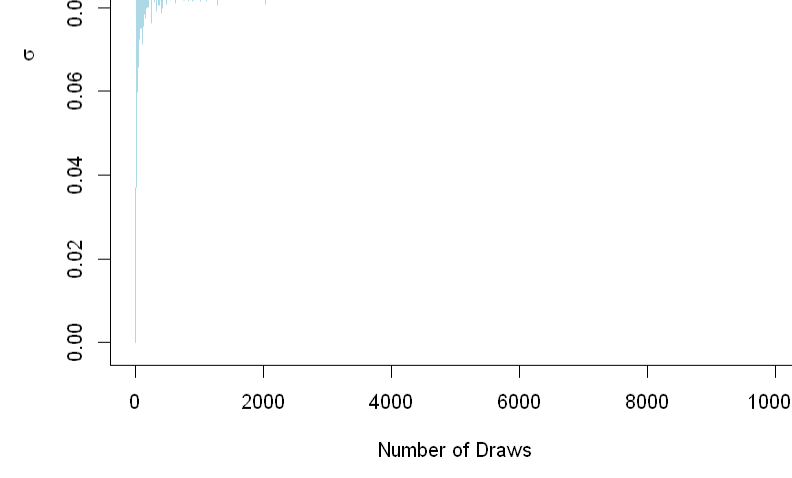
  for(i in 1:n){
    postB[i,] = rbeta(1, alpha, beta)
    res[i,] = c(i, mean(postB), sd(postB))
  }

  return(list(post=postB, mean_sd=res))
}
Posterior = PostB()
```

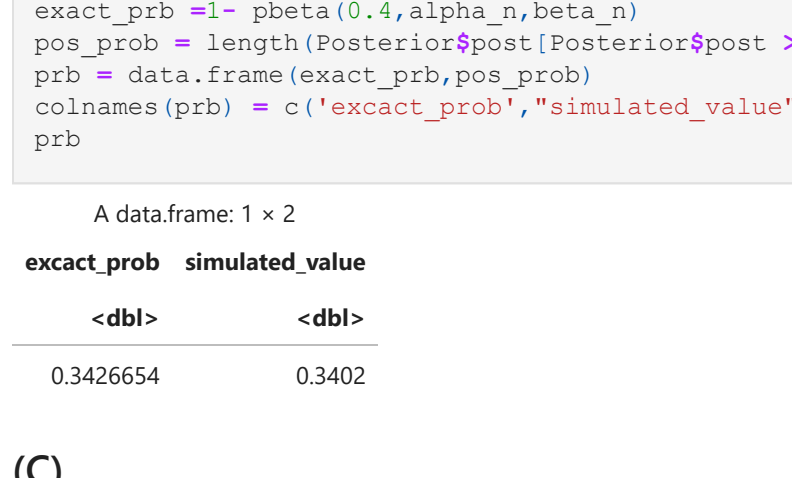
```
In [4]: Posterior$mean_sd[1,3]=0
head(Posterior$mean_sd)
```

A matrix: 6 × 3 of type dbl			
iter	mean	sd	
1	0.3629400	0.0000000	
2	0.2669733	0.07396499	
3	0.3443858	0.06782267	
4	0.3124737	0.09665964	
5	0.4286403	0.13712276	
6	0.3704016	0.11454616	


```
In [5]: x <- seq( from = -5, to = 10 , by = 0.001)
hist(Posterior$post, col = 'lightyellow', freq = FALSE, breaks = 50, xlab = 'posterior', main = 'Histogram of Posterior')
curve(dbeta(x, alpha_n, beta_n), add = TRUE, col = 'blue', lwd = 2, lty = 3)
```



```
In [6]: plot(x=1:NDraws, y=Posterior$mean_sd[,2], col = 'lightblue', lwd = 1, type = 'l', xlab = 'Number of Draws', ylab = 'mu',
  abline(h=TrueMean, col = 'deeppink3', lwd = 3))
```



```
In [7]: plot(x=1:NDraws, y=Posterior$mean_sd[,3], col = 'lightblue', lwd = 1, type = 'l', xlab = 'Number of Draws', ylab = 'sigma',
  abline(h=TrueVar, col = 'deeppink3', lwd = 3))
```



As we can see from the two plots above, as the number of draws increases, both μ and σ converge to the true values.

(b)

we can see that our estimation of $\Pr(\theta > 0.4 | y)$ is very close to true theoretical value.

```
In [8]: exact_prb = 1 - pbeta(0.4, alpha_n, beta_n)
pos_prob = length(Posterior$post[Posterior$post > 0.4]) / length(Posterior$post)
prb = data.frame(exact_prb, pos_prob)
colnames(prb) = c("exact_prob", "simulated_value")
prb
```

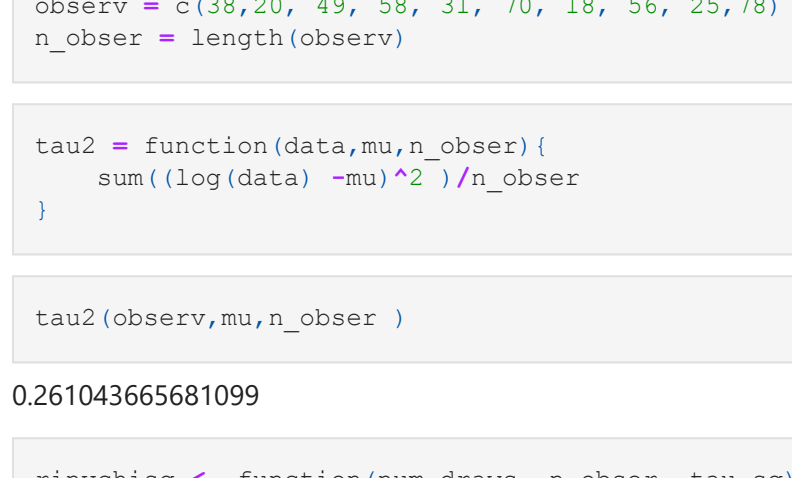
A data frame: 1 × 2

exact_prob	simulated_value
<dbl>	<dbl>
0.3426654	0.3402

(C)

Histogram and kernel density of the data simulated from the posterior distribution of the $\phi = \log(\theta(1 - \theta))$ with 10000 draws.

```
In [9]: phi = log(Posterior$post / (1 - Posterior$post))
hist(phi, breaks = 100, probability = TRUE, col = 'cyan', xlab = expression(phi))
lines(density(phi), col = 'deeppink3', lty = 2, lwd = 3)
```



Assignment 2

Log-normal distribution and the Gini coefficient.

Assume that you have asked 10 randomly selected persons about their incomes (in thousands Swedish Krona) and obtained the following ten observations: 14, 25, 45, 25, 30, 33, 19, 50, 34 and 67. A common model for non-negative continuous variables is the log-normal distribution. The log-normal distribution $\log \mathcal{N}(\mu, \sigma^2)$

has density function

$$p(y | \mu, \sigma^2) = \frac{1}{y \cdot \sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(\log y - \mu)^2\right)$$

where $y > 0$, $\mu > 0$ and $\sigma^2 > 0$

The log-normal distribution is related to the normal distribution as follows: $y \sim \log \mathcal{N}(\mu, \sigma^2)$ then

$\log y \sim \mathcal{N}(\mu, \sigma^2)$ Let $y_1, \dots, y_n | \mu, \sigma^2 \stackrel{iid}{\sim} \log \mathcal{N}(\mu, \sigma^2)$, where $\mu = 3.8$

is assumed to be known but σ^2 is unknown with non-informative prior $p(\sigma^2) \propto 1/\sigma^2$ The posterior for σ^2 is the $\text{Inv} - \chi^2(n, \tau^2)$

distribution, where

$$\tau^2 = \frac{\sum_{i=1}^n (\log y_i - \mu)^2}{n}$$

(a)

Simulate 10, 000 draws from the posterior of σ^2 (assuming $\mu = 3.8$) and compare it with the theoretical $\text{Inv} - \chi^2(n, \tau^2)$ posterior distribution.

```
In [10]: num_draws = 10000 #sample size
mu = 3.8
observ = c(38, 20, 49, 58, 31, 70, 18, 56, 25, 78)
n_observ = length(observ)
```

```
In [11]: tau2 = function(data, mu, n_observ) {
  sum((log(data) - mu)^2) / n_observ
}
```

```
In [12]: tau2(observ, mu, n_observ )
```

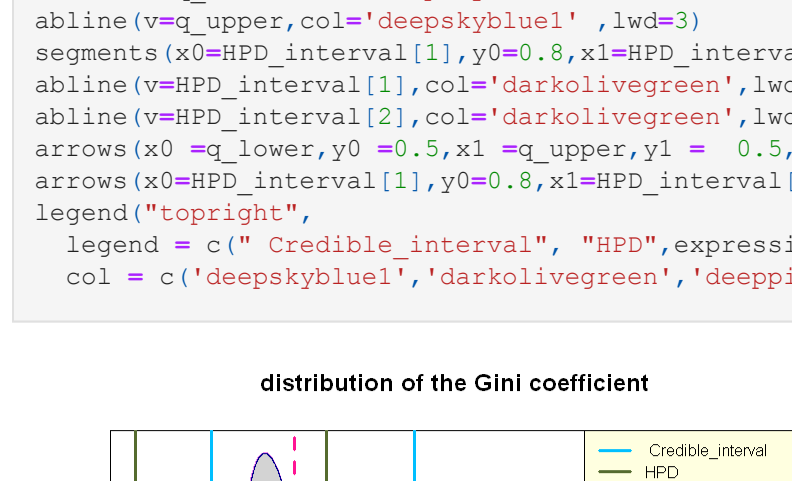
0.261043665681099

```
In [13]: rinvcchi <- function(num_draws, n_observ, tau_sq) {
  set.seed(1234)
  x <- rchisq(num_draws, df = n_observ - 1)
  x_inv <- ((n_observ - 1) * tau_sq) / x
  return(x_inv)
}
```

```
In [14]: post_sigma2 <- function(m) {
  set.seed(12345)
  rinvcchiq(num_draws = num_draws, n_observ = n_observ, tau_sq = tau2(observ, mu, n_observ))
}
```

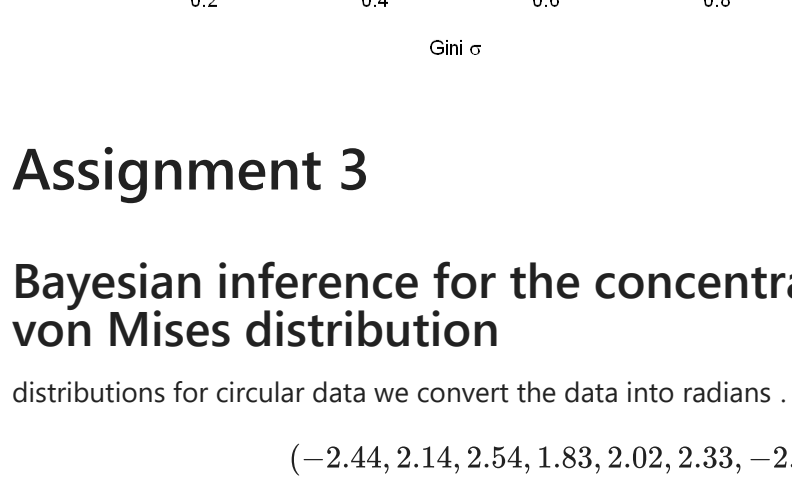
```
In [15]: dinvcchi <- function(x, n_observ, tau_sq) {
  res <- (((tau_sq * (n_observ - 1)) / 2) * (n_observ - 1) / 2 * exp(-(tau_sq * (n_observ - 1)) / (2 * x))) /
  return(res)
}
```

```
In [16]: x <- seq(from = 0, to = 10, by = 0.001)
hist(post_sigma2(m = num_draws), probability = TRUE, col = 'pink', breaks = 100,
  main = expression(paste("Simulated and theoretical posterior of ", sigma^2)),
  xlab = expression(paste("sigma^2")),
  curve(dinvcchi(x, n_observ = n_observ, tau_sq = tau2(observ, mu, n_observ)), add = TRUE, col = 'blue', lwd = 2, lty = 3))
```



The most common measure of income inequality is the Gini coefficient, G , where $0 < G < 1$. $G = 0$ means a completely equal income distribution, whereas $G = 1$ means complete income inequality (see Wikipedia for more information about the Gini coefficient). It can be shown that $G = 2\Phi\left(\sigma/\sqrt{2}\right) - 1$ when incomes follow a $\log \mathcal{N}(\mu, \sigma^2)$ distribution. $\Phi(z)$ is the cumulative distribution function (CDF) for the standard normal distribution with mean zero and unit variance. Use the posterior draws in a) to compute the posterior distribution of the Gini coefficient G for the current data set

```
In [17]: sigma2 = post_sigma2(m)
g_sigma = sqrt(sigma2) / sqrt(2)
gpdf <- 2 * pnorm(q = g_sigma, mean = 0, sd = 1) - 1
hist(gpdl, probability = TRUE, col = 'yellow', breaks = 50, main = 'the posterior distribution of the Gini coefficient', xlab = expression(paste("Gini ", sigma)))
lines(density(gpdl), col = 'darkblue', lty = 2, lwd = 2)
```



(c)

Use the posterior draws from b) to compute a 90% equal tail credible interval for G . A 90% equal tail interval (a: b) cuts off 5% percent of the posterior probability mass to the left of a, and 5% to the right of b. Also, do a kernel density estimate of the posterior of G using the density function in R with default settings, and use that kernel density estimate to compute a 90% Highest Posterior Density Interval (HPDI) for G . Compare the two intervals.

```
In [18]: alpha_conf = 0.1
q_lower = quantile(gpdl, alpha_conf/2)
q_upper = quantile(gpdl, 1 - alpha_conf/2)
c(q_lower, q_upper)
true_mean = mean(gpdl)
```

5%: 0.208822072375841 95%: 0.445627779268848

```
In [19]: true_mean
```

0.305486649437337

HPD: shortest possible interval that under the posterior has the significance probability (ie. 0.9)

```
In [20]: KerDenG <- density(gpdl)
xx = KerDenG$x
yy = sort(KerDenG$y, decreasing = TRUE, index.return = TRUE)
```

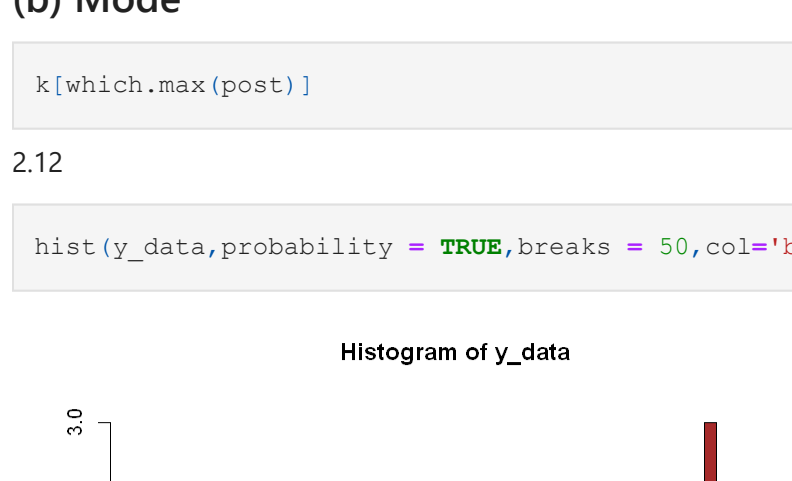
we will use the cumulative sum of kernel density values to find the shortest interval in data that contains 0.9 of the data. the reason we can use cumulative to find HPD interval is because our data is unimodal.

```
In [21]: CumSum = cumsum(yy$x)
HPD_D = CumSum[length(CumSum)] * 0.9
temp = which(CumSum < HPD_D)
HPD_interval = range(xx[which(CumSum < HPD_D)])
```

```
In [22]: df = data.frame('HPD_interval' = HPD_interval, 'credible_interval' = c(q_lower, q_upper))
```

```
In [24]: y_temp = density(gpdl, n = 10000)
plot(x = y_temp$x, y = y_temp$y, type = 'l', lwd = 2, col = 'violet', main = 'distribution of Gini coefficient', xlab = expression(paste("Gini ", sigma)), ylab = 'Density')
polygon(x = y_temp$x, y = y_temp$y, col = 'lightgrey', border = 'darkblue')
abline(v = true_mean, col = 'deeppink', lwd = 3, lty = 2)

segments(x0 = q_lower, y0 = 0.5, x1 = q_upper, y1 = 0.5, col = 'deebskyblue1', lwd = 3)
abline(v = q_lower, col = 'deebskyblue1', lwd = 3)
abline(v = q_upper, col = 'deebskyblue1', lwd = 3)
abline(v = HPD_interval[1], col = 'darkolivegreen', lwd = 3)
abline(v = HPD_interval[2], col = 'darkolivegreen', lwd = 3)
arrows(x0 = q_lower, y0 = 0.5, x1 = q_upper, y1 = 0.5, col = 'deebskyblue1', lwd = 3, code = 3)
arrows(x0 = HPD_interval[1], y0 = 0.8, x1 = HPD_interval[2], y0 = 0.8, col = 'darkolivegreen', lwd = 3)
legend("topright",
  legend = c("Credible_interval", "HPD", expression(mu)),
  col = c("deebskyblue1", "darkolivegreen", 'deeppink'), lty = c(1, 1, 2), cex = 0.8, lwd = 3, bty = "n")
```



Assignment 3

Bayesian inference for the concentration parameter in the von Mises distribution

distributions for circular data we convert the data into radians. The 10 observations in radians are

(-2.44, 2.14, 2.54, 1.83, 2.02, 2.33, -2.79, 2.23, 2.07, 2.02).

Assume that these data points are independent observations following the von Mises distribution

$$p(y | \mu, \kappa) = \frac{\exp[\kappa \cdot \cos(y - \mu)]}{2\pi I_0(\kappa)}, \quad -\pi \leq y \leq \pi,$$

where $I_0(\kappa)$ is the modified Bessel function of the first kind of order zero. The parameter $-\pi \leq \mu \leq \pi$ is the mean direction and $\kappa > 0$ is called the concentration parameter. Large κ gives a small variance around μ , and vice versa. Assume that μ is known to be 2.39. Let $\kappa \sim \text{Exponential}(\lambda = 1)$ a priori, where λ is the rate parameter of the exponential distribution (so that the mean is $1/\lambda$).

(a) Plot the posterior distribution of κ for the wind direction data over a fine grid of κ values.

(b) Find the (approximate) posterior mode of κ from the information in a).

(a) Posterior :

$$p(\kappa | y_1, y_2, \dots, y_n) \propto p(y_1, y_2, \dots, y_n | \kappa) \cdot p(\kappa)$$

$$p(\kappa | y_1, y_2, \dots, y_n) \propto \left[\frac{1}{I_0(\kappa)} \right]^n \cdot \exp \left[\sum_{i=1}^n \kappa \cdot \cos(y_i - \mu) - \lambda \kappa \right]$$

$$p(\kappa | y_1, y_2, \dots, y_n) \propto \left[\frac{1}{I_0(\kappa)} \right]^n \cdot \exp \left[\sum_{i=1}^n \kappa \cdot \cos(y_i - 2.39) - \kappa \right]$$

```
In [25]: y_data = c(-2.44, 2.14, 2.54, 1.83, 2.02, 2.33, -2.79, 2.23, 2.07, 2.02)
```

```
In [26]: n = length(y_data)
mu = 2.39
n
```

10

```
In [27]: # a function to compute prior for kappa (exponential with lambda = 1)
prior_k <- function(k) {
  dexp(x = k, rate = 1)
}
```

```
In [28]: likelihood_k <- function(y, mu, n, k) {
  exp(k * sum(cos(y - mu))) / ((2 * pi)^n * besselI(k, n, n = 0)) ^ n
}
```

```
In [29]: k = seq(from = 0, to = 20, by = 0.01)
```

```
In [30]: #
posterior_k <- function(y = y_data, mu = mu, n = n, k = k) {
  likelihood_k(y, mu, n, k) * prior_k(k)
}
```

```
In [31]: post = posterior_k(y = y_data, mu = mu, n = n, k = k)
```

```
In [32]: plot(y = post, x = k, col = 'blue', type = 'l', lwd = 3)
```


(b) Mode

```
In [33]: k[which.max(post)]
```

2.12

```
In [34]: hist(y_data, probability = TRUE, breaks = 50, col = 'brown')
```


from the plot above we can see that $\kappa = 2.12$ is the actual mode of our data.

```
In [35]: y_data[which(y_data < 0)]
```

-2.44 -2.79